

Agent Implementation

This project implements a Value Based method called Deep Q-Networks. DeepMind leveraged Deep Q-Network (DQN) to build Deep Q-Learning algorithm. This algorithm can learn to play Atari games.

Experience Replay is act of sampling batch of tuples from the replay buffer in order to learn more from individual tuples multiple times, recall rate occurrences, and in general make better use of our experiences. and Fixed Q-Targets.

Reinforcement learning is known to be unstable or even to diverge when a nonlinear function approximator such as a neural network is used to represent the action-value (also known as Q) function²⁰. This instability has several causes: the correlations present in the sequence of observations, the fact that small updates to Q may significantly change the policy and therefore change the data distribution, and the correlations between the action-values (Q) and the target values $r + \gamma \max_{a'} Q(s', a')$. We address these instabilities with a novel variant of Q-learning, which uses two key ideas. First, we used a biologically inspired mechanism termed experience replay^{21–23} that randomizes over the data, thereby removing correlations in the observation sequence and smoothing over changes in the data distribution (see below for details). Second, we used an iterative update that adjusts the action-values (Q) towards target values that are only periodically updated, thereby reducing correlations with the target.

The above article is taken from Human-level control through deep reinforcement learning.

Algorithm: Deep Q-learning

Algorithm: Deep Q-Learning

- Initialize replay memory D with capacity N
- Initialize action-value function \hat{q} with random weights \mathbf{w}
- Initialize target action-value weights $\mathbf{w}^- \leftarrow \mathbf{w}$
- for** the episode $e \leftarrow 1$ to M :
 - Initial input frame x_1
 - Prepare initial state: $S \leftarrow \phi(\langle x_1 \rangle)$
 - for** time step $t \leftarrow 1$ to T :

SAMPLE

Choose action A from state S using policy $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}(S, A, \mathbf{w}))$
Take action A , observe reward R , and next input frame x_{t+1}
Prepare next state: $S' \leftarrow \phi(\langle x_{t-2}, x_{t-1}, x_t, x_{t+1} \rangle)$
Store experience tuple (S, A, R, S') in replay memory D
 $S \leftarrow S'$

LEARN

Obtain random minibatch of tuples (s_j, a_j, r_j, s_{j+1}) from D
Set target $y_j = r_j + \gamma \max_a \hat{q}(s_{j+1}, a, \mathbf{w}^-)$
Update: $\Delta \mathbf{w} = \alpha (y_j - \hat{q}(s_j, a_j, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(s_j, a_j, \mathbf{w})$
Every C steps, reset: $\mathbf{w}^- \leftarrow \mathbf{w}$

Code:

Model.py: QNetwork class is implemented. The network will predict the action to perform depending on the environment observed states. The QNetwork will be used by DQN agent and consists of two fully connected layers. The first layer's input size is provided through state_size and the output size is fc1_units. The second fully connect layer's input size is the output size of the first layer. The output size of the second layer is fc2_units. The default fc1_units and fc2_units are 64. The fc2_units is the input size of the last layer.

Below are the dimension of the QNetwork:

state_size (int): Dimension of each state

action_size (int): Dimension of each action

seed (int): Random seed

fc1_units (int): Number of nodes in first hidden layer

fc2_units (int): Number of nodes in second hidden layer

DQN_agent.py: Class Agent is defined. It has functions `__init__`, `step`, `act`, `learn`, and `soft_update`.

- The `__init__` initializes `state_size` and `action_size`. It initializes the QNetwork and ReplayBuffer.
- The `step` function allows the memory of step taken to be saved for future learning and it allows to sample and learn from the tuples.
- The `act` function returns actions for given state as per current policy. It uses Epsilon Greedy action to balance between Exploitation and Exploration.
- The `learn` function updates the value parameters using given batch of experience tuples.
- The `soft_update` function is called by `learn` to update the target network. This is the Fixed Q-Targets.
- The ReplayBuffer is Fixed-size buffer to store experience tuples.

Navigation.ipynb: This notebook allows the training of Agent on the Banana environment.

- It installs the Banana's Brain environment.
- It exams the state and action spaces.
- It takes random actions in the environment.

The DQN agent uses the following parameters:

`BUFFER_SIZE = int(1e5)` # replay buffer size

`BATCH_SIZE = 64` # minibatch size

`GAMMA = 0.99` # discount factor

`TAU = 1e-3` # for soft update of target parameters


`LR = 5e-4` # learning rate

`UPDATE_EVERY = 4` # how often to update the network

The input size is 37 and consists of two fully connected layers.

The Neural Networks use the Adam optimizer with a learning rate $LR=5e-4$ and are trained using a `BATCH_SIZE=64`

Episode 100 Average Score: 0.60
Episode 200 Average Score: 4.23
Episode 300 Average Score: 7.25
Episode 400 Average Score: 10.08
Episode 500 Average Score: 12.02
Episode 533 Average Score: 13.03
Environment solved in 433 episodes! Average Score: 13.03



Double DQN and Dueling DQN might be able to perform better this Deep Q-Learning.